

Adaptive Cyber-Physical Security: Phase-3 Project Report — Hybrid Learned-Representation Intrusion Detection with Meta-Ensemble Fusion

Pugazhendhi J
Department of CS and AI
Rishihood University
Sonipat, India

pugazhendhi.j23csai@nst.rishihood.edu.in

Dally R
Department of CS and AI
Rishihood University
Sonipat, India

dally.r23csai@nst.rishihood.edu.in

Abstract—Phase 2 of this project demonstrated that combining a supervised Random Forest with anomaly detectors (One-Class SVM and Autoencoder) under OR-logic fusion substantially improves zero-day recall at manageable precision cost. Phase 3 extends these results by constructing six hybrid models that couple learned Autoencoder (AE) bottleneck representations with explicit boundary classifiers—One-Class SVM (OC-SVM), Isolation Forest (IF), and pseudo-label Random Forest (RF)—operating in the AE’s 32-dimensional latent space. Three AE variants are evaluated: a standard AE, a Sparse AE with L1 activity regularization, and a Denoising AE with Gaussian input corruption. A Meta-Ensemble averaging all eight component anomaly scores is also evaluated. Experiments on the full CIC-IDS-2018 dataset (3.2M processed flows, 13 attack families) show that the standalone Denoising AE achieves the highest per-threshold precision (~ 0.82) while the standard AE baseline leads in ROC-AUC (0.8965). The Meta-Ensemble, combining all component scores with equal weight, attains ROC-AUC of 0.4140—highlighting that naive score-level fusion degrades discriminative power when constituent boundary models (OC-SVM, IF) exhibit poor AUC in the compressed latent space. The central finding is that *representation quality* is the bottleneck: AE variants alone consistently outperform their hybrid counterparts when boundary models cannot reliably separate the learned latent distributions.

Index Terms—Intrusion Detection, CIC-IDS-2018, Autoencoder, Sparse Autoencoder, Denoising Autoencoder, One-Class SVM, Isolation Forest, Random Forest, Hybrid IDS, Meta-Ensemble, Zero-Day Detection, Anomaly Detection, Cyber-Physical Security, Deep Learning

I. INTRODUCTION

Modern cyber-physical systems (CPS) face a continuously evolving threat landscape in which both known and previously unseen (zero-day) attacks must be detected with minimal false alarms. Phases 1 and 2 of this project established two key findings on the CIC-IDS-2018 benchmark [1]: (i) a supervised Random Forest (RF) achieves near-perfect precision on known attack families but zero recall on zero-day traffic [2]; and (ii) anomaly detectors such as the One-Class SVM (OC-SVM) [3] and a deep Autoencoder (AE) [4], [5] generalize beyond

the training distribution, enabling non-zero zero-day recall at the cost of lower precision.

Phase 3 investigates whether *learned latent representations* from AE variants can serve as a higher-quality feature space for boundary classifiers, thereby improving the precision–recall trade-off beyond what either paradigm achieves alone. Six hybrid architectures and a Meta-Ensemble are evaluated:

- 1) **AE + OC-SVM**: OC-SVM trained on the standard AE’s 32-dim bottleneck features (benign only).
- 2) **AE + Isolation Forest**: IF [6] trained on the same bottleneck.
- 3) **AE + RF (Pseudo-label)**: RF trained on AE-generated pseudo-labels using the original 68-dim feature space.
- 4) **Sparse AE + OC-SVM**: OC-SVM on L1-regularized bottleneck features.
- 5) **Denoising AE + OC-SVM**: OC-SVM on noise-robust bottleneck features.
- 6) **Meta-Ensemble**: Equal-weight average of all eight normalized anomaly scores [7].

II. BACKGROUND

A. Phase 1 and Phase 2 Summary

Phase 1 trained an RF and an OC-SVM on a subset of CIC-IDS-2018 with a manually configured seen/unseen attack split. The RF delivered perfect precision (1.000) but zero zero-day recall; the OC-SVM achieved 14.5% zero-day recall with 90.7% precision. Phase 2 scaled to the full 8.2M-flow dataset, introduced a deep AE for continuous anomaly scoring, and fused RF and OC-SVM via OR-logic, raising overall recall to ~ 0.89 and F1 to ~ 0.85 [8], [9].

B. Hybrid IDS Motivation

The core hypothesis of Phase 3 is that the AE’s bottleneck layer learns a compact, low-dimensional representation of normal traffic. Training boundary models *in this latent space* should yield tighter decision boundaries than operating in the

original 68-dim feature space, combining the AE’s representation power with the boundary model’s explicit decision surface [10], [11].

III. METHODOLOGY

A. Dataset and Preprocessing

All experiments reuse the preprocessed CIC-IDS-2018 cache from Phase 2: 3,215,110 scaled flows (2,050,783 benign, 1,164,327 attack), 68 features after correlation-based pruning, robust scaling, and zero-value imputation.

B. Autoencoder Architecture

Three AE variants share a common encoder–decoder skeleton: $68 \rightarrow 128 \rightarrow 64 \rightarrow \mathbf{32} \rightarrow 64 \rightarrow 128 \rightarrow 68$, with batch normalization, ReLU activations, and 20% dropout.

- **Standard AE:** Trained on benign-only flows; MSE loss.
- **Sparse AE:** Identical architecture with L1 activity regularization ($\lambda = 10^{-4}$) on the bottleneck, forcing sparse activations and potentially more separable latent dimensions.
- **Denoising AE:** A `GaussianNoise` ($\sigma = 0.1$) layer on the input during training; targets remain clean, forcing the encoder to learn noise-invariant features [5].

All AEs are trained for up to 50 epochs with early stopping (patience 7) and learning-rate reduction on plateau.

C. Boundary Models

- **OC-SVM** ($\nu = 0.05$, RBF kernel, $\gamma=\text{scale}$): Trained on 80,000 subsampled benign bottleneck features.
- **Isolation Forest** (200 estimators, contamination=0.05): Trained on the full benign bottleneck.
- **RF (Pseudo-label)** (200 estimators, max_depth=12, balanced class weights): Trained on AE-thresholded pseudo-labels using the original 68-dim features (500,000 subsample).

D. Score Fusion

For each hybrid, both the AE reconstruction error and the boundary model’s anomaly score are min-max normalized to $[0, 1]$, then combined as:

$$s_{\text{hybrid}} = \alpha \cdot s_{\text{AE}}^{\text{norm}} + (1 - \alpha) \cdot s_{\text{boundary}}^{\text{norm}}, \quad \alpha = 0.5. \quad (1)$$

The Meta-Ensemble averages all eight normalized component scores with equal weight ($1/N$, $N = 8$).

E. Evaluation Protocol

All models are evaluated on the full mixed dataset using a threshold calibrated at the 99th percentile of benign-only validation errors. Metrics include ROC-AUC, Average Precision (AP), F1, Recall (TPR), Precision, and False Positive Rate (FPR).

IV. RESULTS

A. Model Comparison

Table I presents the comprehensive comparison of all ten model configurations evaluated in Phase 3.

Several key observations emerge from the results:

- The three standalone AE variants (Standard, Sparse, Denoising) consistently achieve the highest ROC-AUC values (> 0.88), confirming that reconstruction error is a strong discriminative signal.
- The Denoising AE alone delivers the best per-threshold F1 (0.1433) and the highest precision (0.8179), indicating that noise-robust training produces more separable reconstruction errors.
- All hybrid models that incorporate OC-SVM or IF in the bottleneck space *degrade* ROC-AUC below 0.50, suggesting that these boundary models fail to reliably separate benign from attack latent representations in the 32-dim space.
- The Meta-Ensemble (Hybrid 6) does not recover performance, confirming that averaging poorly discriminative boundary scores with strong AE scores dilutes overall quality.

B. ROC and Precision-Recall Curves

Figures 1 and 2 present the ROC and Precision-Recall curves for all model configurations.




Fig. 1: ROC curves for all hybrid IDS configurations. The standalone AE variants cluster near the upper-left corner with $\text{AUC} > 0.88$, while hybrid models combining boundary classifiers in the bottleneck space perform near or below the random baseline.

C. Metric Comparison

Figure 3 presents a bar-chart comparison of F1, ROC-AUC, and Recall across all models, visually highlighting the performance gap between standalone AE baselines and their hybrid counterparts.

D. Meta-Ensemble Confusion Matrix

Figure 4 shows the confusion matrix of the Meta-Ensemble at the P99 threshold.

TABLE I: Phase 3 — comprehensive model comparison (P99 threshold)

Model	ROC-AUC	Avg Prec.	F1	Recall	Prec.	Acc.
Baseline: Standard AE	0.8965	0.7578	0.0336	0.0174	0.5012	0.6379
Hybrid 1: AE + OC-SVM	0.4600	0.3137	0.0094	0.0048	0.2127	0.6332
Hybrid 2: AE + Isolation Forest	0.4231	0.3392	0.0116	0.0060	0.2508	0.6336
Hybrid 3: AE + RF (pseudo)	0.7920	0.5593	0.0335	0.0174	0.4976	0.6378
Ref: RF alone (pseudo)	0.3433	0.3176	0.0335	0.0173	0.4978	0.6378
Ref: Sparse AE alone	0.8929	0.7597	0.0336	0.0174	0.4986	0.6378
Hybrid 4: Sparse AE + OC-SVM	0.4206	0.3007	0.0190	0.0098	0.3561	0.6350
Ref: Denoising AE alone	0.8867	0.7675	0.1433	0.0785	0.8179	0.6600
Hybrid 5: DAE + OC-SVM	0.4783	0.3186	0.0057	0.0029	0.1398	0.6325
Hybrid 6: Meta-Ensemble	0.4140	0.3024	0.0098	0.0050	0.2202	0.6332




Fig. 2: Precision-Recall curves for all models. The Denoising AE achieves the highest average precision (0.7675), followed closely by the Sparse AE (0.7597) and Standard AE (0.7578).





Fig. 3: F1 / ROC-AUC / Recall comparison across all models. Standalone AE variants dominate on all three metrics; hybrid models with boundary classifiers in the latent space consistently underperform.


 Fig. 4: Meta-Ensemble confusion matrix (P99 threshold). The ensemble detects only 5,842 of 1,164,327 attack flows (recall ≈ 0.005), while generating 20,690 false positives, yielding a precision of only 0.2202.

E. Per-Attack Detection Comparison

Figure 5 compares per-attack-family detection rates between the hybrid models and the standalone AE baseline.




Fig. 5: Per-attack-family detection rate: hybrid models vs. standalone AE. High-volume DoS/DDoS families are detected more reliably, while low-volume stealthy attacks remain challenging for all configurations [12], [13].

F. Threshold Sensitivity

Figure 6 presents the threshold sensitivity analysis for the hybrid configuration.




Fig. 6: Threshold sensitivity: precision, recall, and F1 as a function of the anomaly-score threshold for the hybrid IDS. Operators can slide the threshold to select an operating point appropriate to their deployment context.

G. RF Feature Importance

Figure 7 shows the top features by importance from the pseudo-label Random Forest (Hybrid 3).

V. ANALYSIS

A. Why Boundary Models Fail in the Bottleneck Space

The most striking result of Phase 3 is that coupling a boundary classifier (OC-SVM, IF) with the AE's learned




Fig. 7: RF feature importance from the pseudo-label classifier. The most discriminative features correspond to flow duration, packet-length statistics, and inter-arrival-time measures, consistent with known attack signatures in CIC-IDS-2018.

bottleneck consistently *degrades* performance relative to using the AE reconstruction error alone. Two factors explain this:

- 1) **Overlapping latent distributions.** The AE is trained exclusively on benign traffic. Its bottleneck layer learns to represent benign flows compactly, but attack flows—never seen during training—may project into the same latent region. When benign and attack bottleneck vectors overlap, OC-SVM and IF cannot draw meaningful boundaries.
- 2) **OC-SVM scaling limitations.** The quadratic complexity of OC-SVM forces a training-set cap of 80,000 samples, which may inadequately represent the full diversity of the 1.6M benign training vectors in a 32-dimensional space.

B. Score Fusion Dilution Effect

The Meta-Ensemble’s equal-weight averaging demonstrates a classic dilution problem: when half the component scores (the four boundary-model channels) are near-random discriminators ($AUC < 0.50$), averaging them with the four strong AE channels pulls the aggregate score toward randomness. A weighted fusion scheme or a stacking meta-learner that assigns near-zero weight to weak channels would mitigate this effect.

C. Denoising AE as Best Standalone Detector

The Denoising AE’s superior F1 and precision arise from the noise-corruption training protocol: by forcing the encoder to recover clean inputs from noisy observations, the model learns features that are robust to small perturbations. Attack flows, which deviate more substantially from the benign manifold, exhibit disproportionately higher reconstruction errors, widening the benign–attack gap and improving the threshold operating point.

D. Pseudo-Label RF Performance

Hybrid 3 (AE + RF pseudo-label) is the only hybrid that maintains competitive ROC-AUC (0.7920), because the RF

operates in the *original* 68-dim feature space rather than the compressed bottleneck. However, the pseudo-labels themselves are limited by the AE’s threshold quality: only 40,385 of 1,164,327 true attack flows are pseudo-labeled as attacks, creating severe class imbalance that limits the RF’s recall ceiling.

VI. DISCUSSION

A. Operational Insights

The Phase 3 results produce a clear operational recommendation: *standalone AE variants outperform their hybrid counterparts* when boundary models cannot reliably separate latent distributions. The Denoising AE, with its noise-robust training, is the recommended first-line anomaly detector for this dataset, achieving the best per-threshold F1 (0.1433) and precision (0.8179) among all configurations.

For deployment, a two-tier architecture remains appropriate:

- **Tier 1:** A supervised RF for high-confidence detection of known attack families (100% precision).
- **Tier 2:** A Denoising AE for novelty detection, with an operator-tunable threshold to balance recall and false-alarm rate.

B. Precision–Recall Trade-Off

In CPS security, missed attacks (false negatives) are typically more costly than false alarms. The Denoising AE’s operating point—Precision ≈ 0.82 , Recall ≈ 0.08 at P99—provides high-confidence detections. Lowering the threshold increases recall at the expense of precision, which may be acceptable in high-security environments where alert workloads are manageable.

C. Representation Quality as Bottleneck

The central lesson of Phase 3 is that the quality of the learned representation is the limiting factor. If the AE’s latent space does not separate benign and attack distributions, downstream boundary models have no discriminative signal to exploit. Future work should focus on *contrastive* or *semi-supervised* representation learning that explicitly encourages class separation in the latent space.

VII. LIMITATIONS AND FUTURE WORK

A. Current Limitations

- **Static zero-day split.** The seen/unseen partition is not systematically rotated across attack families.
- **No hyperparameter search.** All configurations use heuristic settings; Bayesian optimization may improve results.
- **Equal-weight fusion.** The Meta-Ensemble assigns uniform weight; a learned stacking meta-classifier could suppress weak channels.
- **OC-SVM scalability.** The 80,000-sample training cap limits representativeness; Deep SVDD or kernel approximations should be evaluated.
- **No concept drift modeling.** The pipeline assumes stationary distributions.
- **Latent-space overlap.** The AE does not enforce benign–attack separation; contrastive losses could address this.

B. Recommended Future Work

- 1) **Leave-one-family-out evaluation.** Rotate the held-out attack class to produce robust zero-day recall estimates.
- 2) **Learned fusion weights.** Replace equal-weight averaging with a stacking meta-learner (e.g., logistic regression on component scores).
- 3) **Contrastive AE training.** Add a contrastive or triplet loss to the AE to encourage latent-space separation.
- 4) **Deep SVDD.** Replace OC-SVM with end-to-end trainable Deep Support Vector Data Description.
- 5) **Online / streaming pipeline.** Integrate incremental learning for real-time CPS deployment.
- 6) **Adversarial robustness.** Evaluate adversarial evasion of the AE-based detectors.

VIII. CONCLUSION

Phase 3 of the Adaptive Cyber-Physical Security project systematically evaluates six hybrid IDS architectures that couple learned AE representations with explicit boundary classifiers, along with a Meta-Ensemble that fuses all component scores.

The principal finding is that *learned representations do not automatically yield superior hybrid detectors*: when boundary models (OC-SVM, IF) cannot reliably separate benign and attack latent vectors, the hybrid's ROC-AUC degrades below the standalone AE baseline. The Denoising AE alone achieves the best per-threshold operating point ($F1=0.1433$, $Precision=0.8179$), confirming that noise-robust training produces the most separable reconstruction-error distributions.

The operational recommendation remains consistent with Phase 2: a multi-tier architecture combining a supervised classifier for known threats with a deep anomaly detector for zero-day novelty is the principled, deployable direction for CPS security. Phase 3 adds the nuance that *representation*

quality is the critical bottleneck, and future work should focus on contrastive and semi-supervised approaches that explicitly enforce class separation in the latent space.

REFERENCES

- [1] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, "Toward generating a new intrusion detection dataset and intrusion traffic characterization," in *Proceedings of the 4th International Conference on Information Systems Security and Privacy (ICISSP)*, 2018, pp. 108–116.
- [2] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [3] B. Schölkopf, J. C. Platt, J. Shawe-Taylor, A. J. Smola, and R. C. Williamson, "Estimating the support of a high-dimensional distribution," *Neural Computation*, vol. 13, no. 7, pp. 1443–1471, 2001.
- [4] I. Goodfellow, Y. Bengio, and A. Courville, "Deep learning." MIT Press, 2016. [Online]. Available: <http://www.deeplearningbook.org>
- [5] P. Vincent, H. Larochelle, I. Lajoie, Y. Bengio, and P.-A. Manzagol, "Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion," in *Journal of Machine Learning Research*, vol. 11, 2010, pp. 3371–3408.
- [6] F. T. Liu, K. M. Ting, and Z.-H. Zhou, "Isolation forest," *Proceedings of the Eighth IEEE International Conference on Data Mining*, pp. 413–422, 2008.
- [7] T. G. Dietterich, "Ensemble methods in machine learning," *Lecture Notes in Computer Science*, vol. 1857, pp. 1–15, 2000.
- [8] V. Chandola, A. Banerjee, and V. Kumar, "Anomaly detection: A survey," *ACM Computing Surveys*, vol. 41, no. 3, pp. 15:1–15:58, 2009.
- [9] A. L. Buczak and E. Guven, "A survey of data mining and machine learning methods for cyber security intrusion detection," *IEEE Communications Surveys & Tutorials*, vol. 18, no. 2, pp. 1153–1176, 2016.
- [10] G. E. Hinton and R. R. Salakhutdinov, "Reducing the dimensionality of data with neural networks," *Science*, vol. 313, no. 5786, pp. 504–507, 2006.
- [11] L. Yang *et al.*, "Hybrid deep learning architecture for network anomaly detection in cps environments," *Sensors*, vol. 23, no. 4, 2023. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC9890381/>
- [12] A. Almalawi *et al.*, "An intrusion detection model to detect zero-day attacks in unseen data using machine learning," *PLOS ONE*, vol. 19, no. 10, p. e0308469, 2024. [Online]. Available: <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0308469>
- [13] K. Guo *et al.*, "A network intrusion detection system based on convolutional neural network," *Sensors*, vol. 23, no. 4, p. 2171, 2023. [Online]. Available: <https://www.mdpi.com/1424-8220/23/4/2171>